

04.advanced

March 21, 2018

1 Advanced Python

1.1 Tips

- On all your script :
- shebang `#!/usr/bin/env python3`
- encoding `# -*- coding: utf-8 -*-`
- underscore have a lot of different meanings
- separator in number `10_000` (only python 3.6)
- last result in the interpreter `_`
- I don't care `_ = f()` (dangerous with internationalization)
- weakly private `_something` (won't be imported with `import *`)
- avoid conflict `list_`
- more private (mangled) `__stuff` -> `_ClassName__mangled_stuff`
- magic methods (also mangled) `__init__`
- for internationalization `_()`
- I/O
- **Always** decode in input
- **Always** encode in output
- modules : import involve execution!
Use

```
if __name__ == '__main__':  
    some_computation()
```
- unpacking dans les boucles

- zip:

```
for name, surname in zip(names, surnames):
    ...
```

- enumerate

```
for index, prime in enumerate(primes):
    ...
```

- False tests :

```
False, 0, None, __nonzero__(), __len__()
```

- lambda functions

```
In [1]: square1 = lambda x: x ** 2
```

```
def square2(x):
    return x ** 2
```

```
print(square1(5))
print(square2(5))
```


25

25

- Be carefull with shared references

a is b # a and b are references to the same object

```
a = 1
b = 3
c = 1
d = 'spam'
```



shared references image

```
In [2]: a = 2
        b = a
        print(a, b)
        a = 3
        print(a, b)
        l1 = [1, 2, 3]
        l2 = l1
        l3 = l1[:]
        print(l1, l2, l3)
        l1[1] = 4
        print(l1, l2, l3)
```

2 2
3 2
[1, 2, 3] [1, 2, 3] [1, 2, 3]
[1, 4, 3] [1, 4, 3] [1, 2, 3]

- **never** use a mutable optionnal arg

```
In [3]: def wrong(a_list=[]):  # never use a mutable optionnal argument
        a_list.append(1)
        return a_list
        print(wrong())
        print(wrong())
        print(wrong([2]))
        print(wrong())

        print()
        def good(a_list=None):
            if a_list is None:
                a_list = []
            a_list.append(1)
            return a_list
        print(good())
        print(good())
        print(good([2]))
        print(good())
```

[1]
[1, 1]
[2, 1]
[1, 1, 1]

[1]
[1]
[2, 1]
[1]

- use a shallow copy to modify a list in a loop (or be very carefull)

```
In [4]: from copy import copy
```

```
def good(my_set):
    for value in copy(my_set):
        if 'bert' in value:
            my_set.remove(value)
    return(my_set)

def wrong(my_set):
    for value in my_set:
        if 'bert' in value:
            my_set.remove(value)
    return(my_set)

print("list ok ", good(["einstein", "albert", "bert", "curie"]))
print("set ok ", good({"einstein", "albert", "bert", "curie"}))

print("list Nok ", wrong(["einstein", "albert", "bert", "curie"]))
print("set Nok ", wrong({"einstein", "albert", "bert", "curie"}))
print("END")
```

```
list ok ['einstein', 'curie']
set ok {'curie', 'einstein'}
list Nok ['einstein', 'bert', 'curie']
```

```
-----
RuntimeError                                Traceback (most recent call last)
```

```
<ipython-input-4-92da89fdbec5> in <module>()
    17
    18 print("list Nok ", wrong(["einstein", "albert", "bert", "curie"]))
--> 19 print("set Nok ", wrong({"einstein", "albert", "bert", "curie"}))
    20 print("END")
```

```
<ipython-input-4-92da89fdbec5> in wrong(my_set)
     8
     9 def wrong(my_set):
--> 10     for value in my_set:
    11         if 'bert' in value:
    12             my_set.remove(value)
```

RuntimeError: Set changed size during iteration

- use exceptions
 - try is very fast but except is very slow
 - nice way to get out of multiples loops/functions at the same time
 - Allows you to be sure that an error had been taken care of

```
In [5]: try:
        print("set  Nok ", wrong({"einstein", "albert", "bert", "curie"}))

        except RuntimeError as e:
            print()
            print("Oops, something went wrong:")
            print(e)
            print("Continuing anyway")
            print()

        print("I'm continuing")
```

```
Oups, something went wrong:
Set changed size during iteration
Continuing anyway
```

```
I'm continuing
```

```
In [6]: import sys
        class MyException(Exception):
            pass

        try:
            for i in range(10):
                for j in range(10):
                    if i == j == 5:
                        raise MyException("Found it")
        except MyException as e:
            print("Out of the loop")
            print(e)
            print("Stop")

            # In a script, use a non-zero return code
            # exit(1)

            # In jupyter you can do
            raise

        print("I will never appear")
```

```
Out of the loop
Found it
Stop
```

```
-----

MyException                                Traceback (most recent call last)

<ipython-input-6-4fd9146a2d0c> in <module>()
      7         for j in range(10):
      8             if i == j == 5:
----> 9                 raise MyException("Found it")
     10 except MyException as e:
     11     print("Out of the loop")

MyException: Found it
```

- Use decorator
- debugging, timing, ...

```
In [7]: from functools import wraps
```

```
def PrintAppel(f):
    def before_f():
        new_f.NbAppels += 1
        print("Entering {}".format(f.__name__))

    def after_f():
        print("Exiting {}".format(f.__name__))
        print("This was the call n° {}".format(new_f.NbAppels))

    @wraps(f)
    def new_f(*args, **xargs):
        before_f()
        res = f(*args, **xargs)
        after_f()
        return res

    new_f.NbAppels = 0

    return new_f
```

```
In [8]: @PrintAppel
def a_function(x):
    return 2 * x
```

```
In [9]: a_function(2)
```

```
Entering a_function
```

```
Exiting a_function
```

```
This was the call nr 1
```

```
Out[9]: 4
```

- Make use of classes in order to isolate your work (and your bugs)

```
In [10]: class egg(object):                                # All objects derived from the same object
    """ Full exemple of a class in python """              # shared attribute between all instances
    total_number = 0

    def __init__(self, number=1):                            # constructor
        """ constructor from number """
        self.number = number                                # Good way of defining attributes
        egg.total_number += number

    @classmethod
    def from_recipe(cls, recipe):                            # Alternative constructor
        """ constructor from recipe """
        return cls(recipe["oeufs"])

    def __del__(self):                                       # destructor (rare)
        """ destructor """
        egg.total_number -= self.number

    def __str__(self):                                       # convert your object into printable string
        """ egg to str convertor """
        return "On a total of {} eggs, I own {}".format(egg.total_number, self.number)

    def how_many(self):                                     # a function of the instance
        """ Return the current number of eggs in the recipe """
        return self.number

    @staticmethod
    def how_many_egg():                                       # a function on the class (rare)
        """ Return the total number of eggs for all recipes """
        return egg.total_number

if __name__ == "__main__":

    fried_egg = egg()
    omelette = egg(3)
    recipe_pancake = {"oeufs":2, "lait":0.5, "farine":300}
    pancake = egg.from_recipe(recipe_pancake)
    print("Fried egg      : ", fried_egg)
```

```

print("Omelette      : ", omelette)
print("Pancake       : ", pancake)
print()
print("{:<12} : {:>5} | {}".format("egg",
                                   "NaN",
                                   egg.how_many_egg()))

print("{:<12} : {:>5} | {}".format("fried_egg",
                                   fried_egg.how_many(),
                                   fried_egg.how_many_egg()))

print("{:<12} : {:>5} | {}".format("omelette",
                                   omelette.how_many(), omelette.how_many_egg()))

print("{:<12} : {:>5} | {}".format("pancake",
                                   pancake.how_many(),
                                   pancake.how_many_egg()))

print()
del omelette
print("{:<12} : {:>5} | {}".format("egg",
                                   "NaN",
                                   egg.how_many_egg()))

print("{:<12} : {:>5} | {}".format("fried_egg",
                                   fried_egg.how_many(),
                                   fried_egg.how_many_egg()))

print("{:<12} : {:>5} | {}".format("pancake",
                                   pancake.how_many(),
                                   pancake.how_many_egg()))

del fried_egg
del pancake

print()
help(egg)

```

```

Fried egg      :  On a total of 6 eggs, I own 1
Omelette       :  On a total of 6 eggs, I own 3
Pancake        :  On a total of 6 eggs, I own 2

```

```

egg           :  NaN | 6
fried_egg     :    1 | 6
omelette      :    3 | 6
pancake       :    2 | 6

```

```

egg           :  NaN | 3
fried_egg     :    1 | 3
pancake       :    2 | 3

```

Help on class egg in module __main__:

```

class egg(builtins.object)
|   Full exemple of a class in python

```



```

| Methods defined here:
|
| __del__(self)
|     destructor
|
| __init__(self, number=1)
|     constructor from number
|
| __str__(self)
|     egg to str convertor
|
| how_many(self)
|     Return the current number of eggs in the recipe
|
| -----
| Class methods defined here:
|
| from_recipe(recipe) from builtins.type
|     constructor from recipe
|
| -----
| Static methods defined here:
|
| how_many_egg()
|     Return the total number of eggs for all recipes
|
| -----
| Data descriptors defined here:
|
| __dict__
|     dictionary for instance variables (if defined)
|
| __weakref__
|     list of weak references to the object (if defined)
|
| -----
| Data and other attributes defined here:
|
| total_number = 0

```

- For launching external program :
- If you don't care about the output of the program `python subprocess.check_call(["cmd", "arg1", "arg2"])`
- otherwise (remember to decode) `python data = subprocess.check_output(["cmd", "arg1", "arg2"]).decode('utf-8')`

```
In [11]: import subprocess
import sys
data = subprocess.check_output([sys.executable, "script.py", "Marc"]).decode('utf-8')
print(data)
```

OH HI MARC

1.2 Packaging

- respect PEP (not only for prettyness)
- docstring (auto-documentation)
- All fonctions
- All classes
- All modules (`__init__.py`)
- All files
- type hinting (that's new)
- Almost totally ignored during execution
- mypy (and more and more IDE) are capable of checking consistency
- The typing module allows you to define complex types
- More and more package are compliant with this

```
In [12]: def greeting(name: str) -> str:
var = "Hello" # type: str
# python 3.7 : var = "Hello" : str

return var + " " + name
```

- pytest (unit-testing)
- auto discovery (use tests folders, `test_truc` function, and `TestMachin` classes)
- allow parametrization

```
In [13]: #ONLY for ipython
import ipytest.magics
import pytest
__file__ = '04.advanced.ipynb'
```

```
In [14]: %%run_pytest[clean] -qq
#this was only for ipython

def test_sorted():
    assert sorted([5, 1, 4, 2, 3]) == [1, 2, 3, 4, 5]

# as does parametrize
@pytest.mark.parametrize('input, expected', [
    ([2, 1], [1, 2]),
    ('zasdqw', list('adqswz')),
])
```

```

    )
def test_exemples(input, expected):
    actual = sorted(input)
    assert actual == expected

```

...

- gettext (auto-internationalization) ?
- argparse
- configParser
- logging
- print -> go to console (for ordinary usage)
- warning.warn -> go to console (usually once : for signaling a something the user should fix)
- logging.level -> go anywhere you want (for detailed output and/or diagnostic)

```

In [15]: import logging
import warnings

```

```

def prepare_logging():
    """
    Prepare all logging facilities

    This should be done in a separate module
    """

    # if not already done, initialize logging facilities
    logging.basicConfig()

    # create a logger for the current module
    logger = logging.getLogger(__name__)

    ## ONLY FOR IPYTHON
    # clean logger (ipython + multiple call)
    from copy import copy
    for handler in copy(logger.handlers):
        logger.removeHandler(handler)
    # Do not propagate message to ipython (or else thy will be printed twice)
    logger.propagate=False
    ## ONLY FOR IPYTHON

    # optionnal : change format of the log
    logFormatter = logging.Formatter("%(asctime)s [%(threadName)-12.12s] [%(levelname)-"

    # optionnal : create a handler for file output
    fileHandler = logging.FileHandler("{logPath}/{fileName}.log".format(logPath=".", fi

```

```

# optionnal : create a handler for console output
consoleHandler = logging.StreamHandler()

# optionnal : Apply formatter to both handles
fileHandler.setFormatter(logFormatter)
consoleHandler.setFormatter(logFormatter)

# optionnal : attach handler to the logger
logger.addHandler(fileHandler)
logger.addHandler(consoleHandler)

# what severity to log (default is NOTSET, i.e. all)
logger.setLevel(logging.DEBUG)          # ALL
fileHandler.setLevel(logging.INFO)       # NO DEBUG
consoleHandler.setLevel(logging.WARNING)  # ONLY WARNING AND ERRORS

return logger

def egg():
    warnings.warn("A warning only once")

if __name__ == "__main__":

    logger = prepare_logging()

    egg()

    logger.info('Start reading database')

    records = {'john': 55, 'tom': 66}

    logger.debug('Records: {}'.format(records))
    logger.info('Updating records ...')
    logger.warning("There is only 2 record !")
    logger.info('Saving records ...')
    logger.error("Something happend, impossible to save the records")
    logger.info('Restoring records ...')
    logger.critical("Database corrupted !")
    logger.info('End of program')

    egg()

```

04.advanced.ipynb:53: UserWarning: A warning only once

"outputs": [

2018-03-21 13:43:53,927 [MainThread] [WARNI] There is only 2 record !

2018-03-21 13:43:53,928 [MainThread] [ERROR] Something happend, impossible to save the record

2 Performance

- profiling : Only optimize the bottlenecks !
- timeit (for small snippets of code)

```
In [16]: %timeit [1 + i for i in range(1,10000)]
          %timeit [1 * i for i in range(1,10000)]
          %timeit [1 / i for i in range(1,10000)]
          %timeit [1 // i for i in range(1,10000)]

          %timeit [1. + float(i) for i in range(1,10000)]
          %timeit [1. * float(i) for i in range(1,10000)]
          %timeit [1. / float(i) for i in range(1,10000)]
          %timeit [1. // float(i) for i in range(1,10000)]
```

```
509  $\mu$ s  $\pm$  3.22  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1000 loops each)
465  $\mu$ s  $\pm$  5.18  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1000 loops each)
481  $\mu$ s  $\pm$  1.15  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1000 loops each)
398  $\mu$ s  $\pm$  4.64  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1000 loops each)
1.57 ms  $\pm$  12.7  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1000 loops each)
1.56 ms  $\pm$  5.31  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1000 loops each)
1.59 ms  $\pm$  15.1  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1000 loops each)
2.24 ms  $\pm$  40.4  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 100 loops each)
```

- cProfile (for real code)
python3 -m cProfile -o profile.pstats script.py
gprof2dot -f pstats profile.pstats | dot -Tpng -o profile.png

```
In [17]: import numpy as np
          import cProfile
          import re

          def function2(array):
              for i in range(500):
                  array += 3
                  array = array * 2
              return array

          def function1():
              array = np.random.randint(500000, size=5000000)
              array = function2(array)
              return sorted(array)

          cProfile.run('function1()', sort="tottime")
```

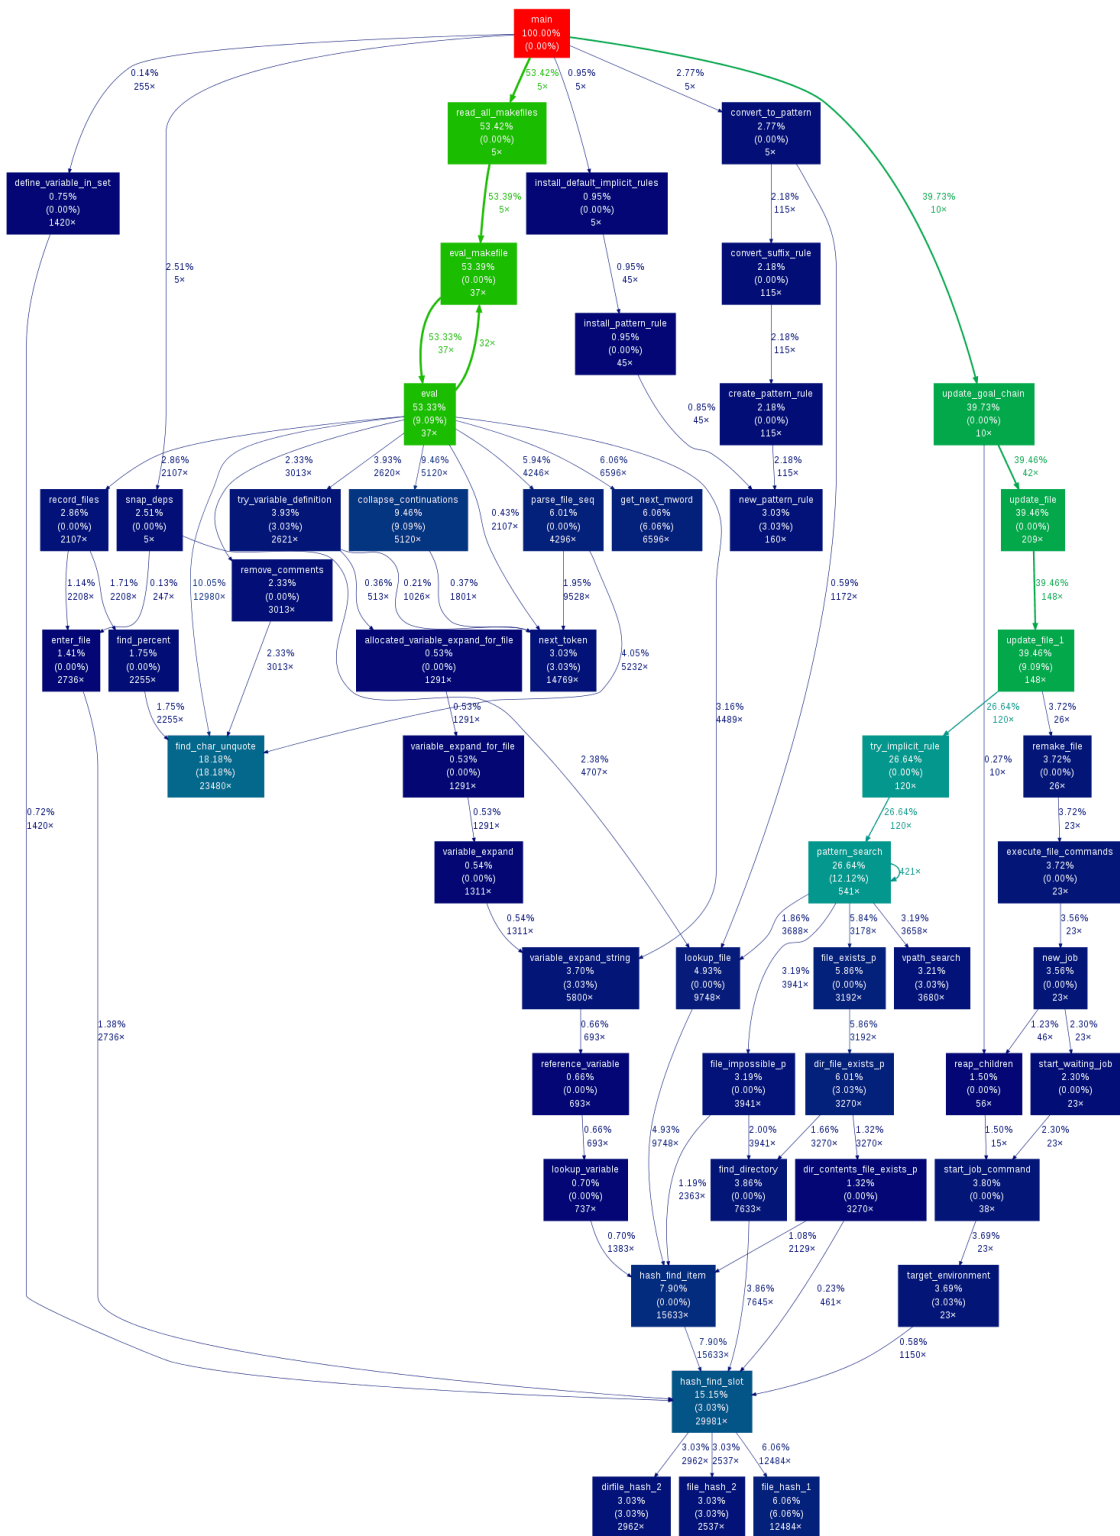
7 function calls in 5.105 seconds

Ordered by: internal time

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	4.786	4.786	4.786	4.786	<ipython-input-17-d5116138ace8>:5(function2)
1	0.251	0.251	0.251	0.251	{built-in method builtins.sorted}
1	0.039	0.039	5.105	5.105	<string>:1(<module>)
1	0.028	0.028	0.028	0.028	{method 'randint' of 'mtrand.RandomState' objects}
1	0.000	0.000	5.066	5.066	<ipython-input-17-d5116138ace8>:11(function1)
1	0.000	0.000	5.105	5.105	{built-in method builtins.exec}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}

```
In [18]: from IPython.display import Image
         Image(filename='images/profile.png')
```

Out[18]:



- in sequential and in Python

- small is beautiful (PEP 20)
 - IO cost a lot (avoid reading and writing into files)
 - choose the right data structure / algorithm
 - prefer numpy based array
 - avoid loops (vectorization using slice)
 - avoid copy of array
 - changing size of an array > Stop optimizing your Python code here (and compile it)
 - inline manually
 - local is faster than global (and avoid dots)
 - use the out argument in numpy
- compiler
 - Almost no modification to your code
 - numexpr (only small expression)
 - numba
 - Almost need a total rewrite
 - cython
 - * compilation must be done separately
 - f2py
 - * included with numpy
 - * compilation must be done separately
 - * be careful to the memory ordering

2.1 Python code (reference)

```
In [19]: import numpy as np

def pyfunction1(a, b):
    return a * b - 4.1 * a > 2.5 * b

def pyfunction2(a, b):
    return np.sin(a) + np.arcsinh(a / b)

def convolve_python(f, g):
    # f is an image and is indexed by (v, w)
    # g is a filter kernel and is indexed by (s, t),
    # it needs odd dimensions
    # h is the output image and is indexed by (x, y),

    if g.shape[0] % 2 != 1 or g.shape[1] % 2 != 1:
        raise ValueError("Only odd dimensions on filter supported")

    # smid and tmid are number of pixels between the center pixel
    # and the edge, ie for a 5x5 filter they will be 2.
    vmax = f.shape[0]
    wmax = f.shape[1]
```



```

smax = g.shape[0]
tmax = g.shape[1]

smid = smax // 2
tmid = tmax // 2

# Allocate result image.
h = np.zeros_like(f)

npsum = np.sum

# Do convolution
for x in range(smid, vmax - smid):
    for y in range(tmid, wmax - tmid):
        # Calculate pixel value for h at (x,y). Sum one component
        # for each pixel (s, t) of the filter g.

        #value = 0
        #for s in range(smax):
        #    for t in range(tmax):
        #        v = x - smid + s
        #        w = y - tmid + t
        #        value += g[s, t] * f[v, w]
        #h[x, y] = value

        # exploit vectorization

        v1 = x - smid
        v2 = v1 + smax
        w1 = y - tmid
        w2 = w1 + tmax
        h[x, y] = npsum(g * f[v1:v2, w1:w2])

    return h

In [20]: def init_copy(f, g):
        return f.copy(), g.copy()

def finish(res):
    return res

def check(fun, data, ref=None, setup=init_copy, finish=finish, timing=True):

    print("Testing ", fun.__name__)

    f, g = setup(*data)

    if ref is None:

```

```

        res = finish(fun(f, g))
    else:
        res = finish(fun(f, g))

    if type(ref) is not type(res):
        print("types are different : ", type(ref), type(res))
        return res

    if ref.dtype != res.dtype:
        print("dtypes are different : ", ref.dtype, res.dtype)
        return res

    if not np.array_equal(res, ref):
        print("results are different")
        print("  ref shape: ", ref.shape)
        print("  res shape: ", res.shape)
        print()
        print("  ref\n", ref)
        print("  res\n", res)
        #print()
        #print("  ref\n", ref[4:-4, 4:-4])
        #print("  res\n", res[4:-4, 4:-4])
        return res

    if timing:
        %timeit fun(f, g)

    return res

```

In [21]: # Create data

```

a = np.arange(1, 1e6)
b = np.arange(1, 1e6)

n = 4
s = 2 * n + 1
g = np.arange(s ** 2, dtype=np.int).reshape((s, s))

N = 200
small_f = np.arange(N * N, dtype=np.int).reshape((N, N))
N = 2000
large_f = np.arange(N * N, dtype=np.int).reshape((N, N))

```

In [22]: # Initial result

```

ref1 = check(pyfunction1, (a, b))
ref2 = check(pyfunction2, (a, b))
print()
ref_small_convolve = check(convolve_python, (small_f, g))

```

Testing pyfunction1

4.84 ms ± 407 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)
Testing pyfunction2
60.1 ms ± 864 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)

Testing convolve_python
167 ms ± 1.17 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)

2.1.1 We can look at the bytecode (halfway to assembly)

And use it to optimize some things

```
In [23]: import dis
         print(dis.code_info(pyfunction2))
         print("Code :")
         dis.dis(pyfunction2)
         print()
```

```
Name:          pyfunction2
Filename:       <ipython-input-19-7d07a1a4a73c>
Argument count: 2
Kw-only arguments: 0
Number of locals: 2
Stack size:    4
Flags:         OPTIMIZED, NEWLOCALS, NOFREE
Constants:
  0: None
Names:
  0: np
  1: sin
  2: arcsinh
Variable names:
  0: a
  1: b
Code :
  7          0 LOAD_GLOBAL              0 (np)
          3 LOAD_ATTR                1 (sin)
          6 LOAD_FAST                 0 (a)
          9 CALL_FUNCTION              1 (1 positional, 0 keyword pair)
         12 LOAD_GLOBAL              0 (np)
         15 LOAD_ATTR                2 (arcsinh)
         18 LOAD_FAST                 0 (a)
         21 LOAD_FAST                 1 (b)
         24 BINARY_TRUE_DIVIDE
         25 CALL_FUNCTION              1 (1 positional, 0 keyword pair)
         28 BINARY_ADD
         29 RETURN_VALUE
```

```
In [24]: from numpy import sin, arcsinh
```

```
# We can avoid the step 0 and 12
def pyfunction2_bis(a, b):
    return sin(a) + arcsinh(a / b)
```

```
In [25]: import dis
print(dis.code_info(pyfunction2_bis))
print("Code :")
dis.dis(pyfunction2_bis)
print()
```

```
Name:          pyfunction2_bis
Filename:      <ipython-input-24-d3f06e8d7add>
Argument count: 2
Kw-only arguments: 0
Number of locals: 2
Stack size:    4
Flags:         OPTIMIZED, NEWLOCALS, NOFREE
Constants:
  0: None
Names:
  0: sin
  1: arcsinh
Variable names:
  0: a
  1: b
Code :
  5          0 LOAD_GLOBAL              0 (sin)
          3 LOAD_FAST                 0 (a)
          6 CALL_FUNCTION              1 (1 positional, 0 keyword pair)
          9 LOAD_GLOBAL              1 (arcsinh)
         12 LOAD_FAST                 0 (a)
         15 LOAD_FAST                 1 (b)
         18 BINARY_TRUE_DIVIDE
         19 CALL_FUNCTION              1 (1 positional, 0 keyword pair)
         22 BINARY_ADD
         23 RETURN_VALUE
```

```
In [26]: _ = check(pyfunction2_bis, (a, b), ref=ref2)
```

Testing pyfunction2_bis

59.8 ms ± 203 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)

2.2 When possible use already available function

```
In [27]: from scipy.signal import convolve2d
```

```
def scipy_convolve(f, g):

    vmax = f.shape[0]
    wmax = f.shape[1]

    smax = g.shape[0]
    tmax = g.shape[1]

    smid = smax // 2
    tmid = tmax // 2

    h = np.zeros_like(f)
    h[smid:vmax - smid, tmid:wmax - tmid] = convolve2d(f, g, mode="valid")

    return h

def scipy_setup(f, g):
    # for some reason, scipy take the filter in reverse...
    gr = g[::-1,::-1]
    return f.copy(), gr.copy()
```

```
In [28]: _ = check(scipy_convolve, (small_f, g), ref=ref_small_convolve, setup=scipy_setup)
        ref_large_convolve = check(scipy_convolve, (large_f, g), setup=scipy_setup)
```

Testing scipy_convolve

6.16 ms ± 57.1 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

Testing scipy_convolve

659 ms ± 6.49 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

2.3 numexpr allows compilation of very simple code

And is multithreaded

```
In [29]: import numexpr as ne
```

```
def ne_function1(a, b):
    return ne.evaluate('a * b - 4.1 * a > 2.5 * b')

def ne_function2(a, b):
    return ne.evaluate("sin(a) + arcsinh(a / b)")
```

```
In [30]: _ = check(ne_function1, (a, b), ref=ref1)
        _ = check(ne_function2, (a, b), ref=ref2)
```

```
Testing ne_function1
977  $\mu$ s  $\pm$  5.74  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1000 loops each)
Testing ne_function2
12.4 ms  $\pm$  616  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 100 loops each)
```

2.4 Numba allows compilation of more complex code using only decorators

And can parallelize part of your code

But it doesn't work everytime

```
In [31]: import numba as nb

@nb.jit(nopython=True, nogil=True, cache=False, parallel=True)
def nb_function1(a, b):
    return a * b - 4.1 * a > 2.5 * b

@nb.jit(nopython=True, nogil=True, cache=False, parallel=True)
def nb_function2(a, b):
    return np.sin(a) + np.arcsinh(a / b)

@nb.jit(nopython=True, nogil=True, cache=False, parallel=True)
def convolve_kernel(f, g):
    # smid and tmid are number of pixels between the center pixel
    # and the edge, ie for a 5x5 filter they will be 2.
    vmax = f.shape[0]
    wmax = f.shape[1]
    smax = g.shape[0]
    tmax = g.shape[1]

    smid = smax // 2
    tmid = tmax // 2

    # Allocate result image.
    h = np.zeros_like(f)

    # Do convolution
    for x in range(smid, vmax - smid):
        for y in range(tmid, wmax - tmid):
            # Calculate pixel value for h at (x,y). Sum one component
            # for each pixel (s, t) of the filter g.

            value = 0
            for s in range(smax):
                for t in range(tmax):
                    v = x - smid + s
                    w = y - tmid + t
                    value += g[s, t] * f[v, w]
```

```

        h[x, y] = value

        #v1 = x - smid
        #v2 = v1 + smax
        #w1 = y - tmid
        #w2 = w1 + tmax
        #h[x, y] = np.sum(g * f[v1:v2, w1:w2])
    return h

@nb.jit(nopython=True, nogil=True, cache=False, parallel=True)
def convolve_numba(f, g):

    if g.shape[0] % 2 != 1 or g.shape[1] % 2 != 1:
        raise ValueError("Only odd dimensions on filter supported")

    return convolve_kernel(f, g)

# Stencil contains implicit loops
@nb.stencil(standard_indexing="g", neighborhood=((-4, 4), (-4, 4)))
def convolve_kernel1(f, g):

    smax = g.shape[0]
    tmax = g.shape[1]

    smid = smax // 2
    tmid = tmax // 2

    h = 0
    for s in range(smax):
        for t in range(tmax):
            h += g[s, t] * f[s - smid, t - tmid]

    return h

@nb.jit(nopython=True, nogil=True, cache=False, parallel=True)
def convolve_numba1(f, g):

    if g.shape[0] % 2 != 1 or g.shape[1] % 2 != 1:
        raise ValueError("Only odd dimensions on filter supported")

    return convolve_kernel1(f, g).astype(f.dtype)

```

```

In [32]: # It will be compiled at first use
         tmpa = np.arange(1, 5)
         tmpb = np.arange(1, 5)
         nb_function1(tmpa, tmpb)
         nb_function2(tmpa, tmpb)

```

```

N = 10
tmpf = np.arange(N * N, dtype=np.int).reshape((N, N))
_ = convolve_numba(tmpf, g)
_ = convolve_numba1(tmpf, g)
print("compilation OK")

```

compilation OK

```

In [33]: _ = check(nb_function1, (a, b), ref=ref1)
_ = check(nb_function2, (a, b), ref=ref2)
print()
_ = check(convolve_numba, (small_f, g), ref=ref_small_convolve)
_ = check(convolve_numba, (large_f, g), ref=ref_large_convolve)
print()
_ = check(convolve_numba1, (small_f, g), ref=ref_small_convolve)
_ = check(convolve_numba1, (large_f, g), ref=ref_large_convolve)

```

Testing nb_function1

797 μ s \pm 8.59 μ s per loop (mean \pm std. dev. of 7 runs, 1000 loops each)

Testing nb_function2

14.1 ms \pm 169 μ s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

Testing convolve_numba

2.35 ms \pm 23.6 μ s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

Testing convolve_numba

251 ms \pm 2.42 ms per loop (mean \pm std. dev. of 7 runs, 1 loop each)

Testing convolve_numba1

1.51 ms \pm 10.9 μ s per loop (mean \pm std. dev. of 7 runs, 1000 loops each)

Testing convolve_numba1

101 ms \pm 1.04 ms per loop (mean \pm std. dev. of 7 runs, 10 loops each)

2.5 Cython is the most efficient way to optimize your code

But you have to: - type *every* variable - explicit all loops - parallelize manually - compile it separately (or use ipython magic) - dependencies - Windows : - Visual studio build tools - or Mingw - or Windows Subsystem for Linux - Linux : - gcc - or icc

```

In [34]: %load_ext Cython

```

```

In [35]: %%cython -a

```

```

# cython: language_level=3
# cython: initializedcheck=False
# cython: binding=True
# cython: nonecheck=False
# distutils: extra_compile_args = -fopenmp -march=native -Ofast -fvector-cost-model=cheap
# distutils: extra_link_args = -fopenmp

```



```

### cython: np_pythran=True
import numpy as np
cimport numpy as np
cimport cython
from libc.math cimport sin
from libc.math cimport asinh
from cython.parallel cimport parallel, prange

@cython.boundscheck(False) # turn off bounds-checking for entire function
@cython.wraparound(False) # turn off negative index wrapping for entire function
def cfunction1(double[:,::1] a, double[:,::1] b):
    cdef long n = a.shape[0]

    cdef long[:] res = np.empty([n], dtype=long)

    cdef Py_ssize_t i
    for i in prange(n, nogil=True, num_threads=8):
        res[i] = a[i] * b[i] - 4.1 * a[i] > 2.5 * b[i]
    return np.asarray(res, dtype=bool)

@cython.boundscheck(False) # turn off bounds-checking for entire function
@cython.wraparound(False) # turn off negative index wrapping for entire function
def cfunction2(double[:,::1] a, double[:,::1] b):
    cdef long n = a.shape[0]

    cdef double[:] res = np.empty([n], dtype=np.double)

    cdef Py_ssize_t i
    for i in prange(n, nogil=True, num_threads=8):
        res[i] = sin(a[i]) + asinh(a[i] / b[i])
    return np.asarray(res)

@cython.boundscheck(False) # turn off bounds-checking for entire function
@cython.wraparound(False) # turn off negative index wrapping for entire function
def convolve_cython(long[:,::1] f, long[:,::1] g):
    cdef long vmax = f.shape[0]
    cdef long wmax = f.shape[1]
    cdef long smax = g.shape[0]
    cdef long tmax = g.shape[1]

    # f is an image and is indexed by (v, w)
    # g is a filter kernel and is indexed by (s, t),
    # it needs odd dimensions
    # h is the output image and is indexed by (x, y),
    if smax % 2 != 1 or tmax % 2 != 1:
        raise ValueError("Only odd dimensions on filter supported")

```

```

# smid and tmid are number of pixels between the center pixel
# and the edge, ie for a 5x5 filter they will be 2.

cdef long smid = smax // 2
cdef long tmid = tmax // 2

# Allocate result image.
cdef long[:,::1] h = np.zeros([vmax, wmax], dtype=long)

cdef long value
cdef Py_ssize_t x, y, s, t, v, w

# Do convolution
for x in prange(smid, vmax - smid, nogil=True, num_threads=8):
    for y in range(tmid, wmax - tmid):
        # Calculate pixel value for h at (x,y). Sum one component
        # for each pixel (s, t) of the filter g.

        value = 0
        for s in range(smax):
            for t in range(tmax):
                v = x - smid + s
                w = y - tmid + t
                value = value + g[s, t] * f[v, w]
        h[x, y] = value

    return np.asarray(h)

```

Out[35]: <IPython.core.display.HTML object>

```

In [36]: _ = check(cfunction1, (a, b), ref=ref1)
         _ = check(cfunction2, (a, b), ref=ref2)
         print()
         _ = check(convolve_cython, (small_f, g), ref=ref_small_convolve)
         _ = check(convolve_cython, (large_f, g), ref=ref_large_convolve)

```

Testing cfunction1

2.1 ms ± 296 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

Testing cfunction2

15 ms ± 126 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

Testing convolve_cython

655 µs ± 6.15 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)

Testing convolve_cython

68.8 ms ± 608 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)

```

In [37]: csource = """
         #include "CModule.h"

```

```

void convolve_c (const long f[],
                 const long g[],
                 long h[],
                 const size_t vmax,
                 const size_t wmax,
                 const size_t smax,
                 const size_t tmax)
{
    const size_t smid = smax / 2;
    const size_t tmid = tmax / 2;

    for(size_t s = 0; s < vmax * wmax; ++s) {
        h[s] = 0;
    }

    // Do convolution
    #pragma omp parallel for default(shared) num_threads(8)
    for(size_t x = smid; x < vmax - smid; ++x) {
        for(size_t y = tmid; y < wmax - tmid; ++y) {
            // Calculate pixel value for h at (x,y).
            // Sum one component for each pixel (s, t) of the filter g.

            long value = 0;
            for(size_t s = 0; s < smax; ++s) {
                for(size_t t = 0; t < tmax; ++t) {
                    size_t v = x - smid + s;
                    size_t w = y - tmid + t;
                    value = value + g[s*tmax + t] * f[v*wmax + w];
                }
            }
            h[x*wmax + y] = value;
        }
    }
}

"""

hsource = """
#include <stddef.h>

void convolve_c (const long f[],
                 const long g[],
                 long h[],
                 const size_t vmax,
                 const size_t wmax,
                 const size_t smax,
                 const size_t tmax);

```

```

    """

with open("CModule.c", 'w') as cfile:
    for line in csource:
        cfile.write(line)

with open("CModule.h", 'w') as hfile:
    for line in hsource:
        hfile.write(line)

In [38]: %%cython -a
# cython: language_level=3
# cython: initializedcheck=False
# cython: binding=True
# cython: nonecheck=False
# distutils: sources = CModule.c
# distutils: extra_compile_args = -fopenmp -march=native -Ofast -fvect-cost-model=cheap
# distutils: extra_link_args = -fopenmp
### cython: np_pythran=True
import numpy as np
cimport numpy as np
cimport cython

cdef extern from "CModule.h":
    long* convolve_c (const long f[],
                      const long g[],
                      long h[],
                      const size_t vmax,
                      const size_t wmax,
                      const size_t smax,
                      const size_t tmax) nogil

@cython.boundscheck(False) # turn off bounds-checking for entire function
@cython.wraparound(False) # turn off negative index wrapping for entire function
def convolve_cython2(long[:, ::1]& f, long[:, ::1]& g):
    # f is an image and is indexed by (v, w)
    # g is a filter kernel and is indexed by (s, t),
    # it needs odd dimensions
    # h is the output image and is indexed by (x, y),

    cdef long vmax = f.shape[0]
    cdef long wmax = f.shape[1]
    cdef long smax = g.shape[0]
    cdef long tmax = g.shape[1]

    if smax % 2 != 1 or tmax % 2 != 1:
        raise ValueError("Only odd dimensions on filter supported")

```

```

cdef long[:,::1] h = np.empty([vmax, wmax], dtype=long)

# Do convolution
with nogil:
    convolve_c(&f[0,0],
               &g[0,0],
               &h[0,0], vmax, wmax, smax, tmax)

return np.asarray(h)

```

Out[38]: <IPython.core.display.HTML object>

```

In [39]: _ = check(convolve_cython2, (small_f, g), ref=ref_small_convolve)
         _ = check(convolve_cython2, (large_f, g), ref=ref_large_convolve)

```

Testing convolve_cython2

566 μ s \pm 19.1 μ s per loop (mean \pm std. dev. of 7 runs, 1000 loops each)

Testing convolve_cython2

56.4 ms \pm 509 μ s per loop (mean \pm std. dev. of 7 runs, 10 loops each)

2.6 Fortran through f2py is also very efficient

But you have to - rewrite your code - compile it separately

```

In [40]: fsource = """
        module fortranmodule
        use OMP_LIB
        implicit none
        contains

        subroutine function1(a, b, c, n)
        implicit none
        integer(kind=8), intent(in) :: n
        double precision,intent(in) :: a(n)
        double precision,intent(in) :: b(n)

        logical,intent(out)          :: c(n)

        !$OMP PARALLEL WORKSHARE NUM_THREADS(8)
        c = a * b - 4.1 * a > 2.5 * b
        !$OMP END PARALLEL WORKSHARE

        end subroutine function1

        subroutine function2(a, b, c, n)
        implicit none

```

```

integer(kind=8), intent(in)  :: n
double precision,intent(in)  :: a(n)
double precision,intent(in)  :: b(n)

double precision,intent(out) :: c(n)

!$OMP PARALLEL WORKSHARE NUM_THREADS(8)
c = sin(a) + asinh(a / b)
!$OMP END PARALLEL WORKSHARE

end subroutine function2

subroutine convolve_fortran(f, g, vmax, wmax, smax, tmax, h, err)
implicit none
integer(kind=8),intent(in)  :: vmax,wmax,smax,tmax
integer(kind=8),intent(in)  :: f(vmax, wmax), g(smax, tmax)

integer(kind=8),intent(out) :: h(vmax, wmax)
integer(kind=8),intent(out) :: err

integer(kind=8) :: smid,tmid
!integer(kind=8) :: value
integer(kind=8) :: x, y
!integer(kind=8) :: v, w, s, t
integer(kind=8) :: v1,v2,w1,w2

! f is an image and is indexed by (v, w)
! g is a filter kernel and is indexed by (s, t),
!   it needs odd dimensions
! h is the output image and is indexed by (v, w),

err = 0
if (modulo(smax, 2) /= 1 .or. modulo(tmax, 2) /= 1) then
    err = 1
    return
endif

! smid and tmid are number of pixels between the center pixel
! and the edge, ie for a 5x5 filter they will be 2.
smid = smax / 2
tmid = tmax / 2

h = 0
! Do convolution
! warning : memory layout is different in fortran
! warning : array start at 1 in fortran

```

```

!$OMP PARALLEL DO DEFAULT(SHARED) COLLAPSE(1) &
!$OMP PRIVATE(v1,v2,w1,w2) NUM_THREADS(8)
do y = tmid + 1, wmax - tmid
    do x = smid + 1, vmax - smid
        ! Calculate pixel value for h at (x,y). Sum one component
        ! for each pixel (s, t) of the filter g.

        !value = 0
        !do t = 1, tmax
        !    do s = 1, smax
        !        v = x - smid + s - 1
        !        w = y - tmid + t - 1
        !
        !        value = value + g(s, t) * f(v, w)
        !    enddo
        !enddo

        v1 = x - smid
        v2 = v1 + smax
        w1 = y - tmid
        w2 = w1 + tmax
        h(x, y) = sum(g(1:smax,1:tmax) * f(v1:v2,w1:w2))
    enddo
enddo
!$OMP END PARALLEL DO
return
end subroutine convolve_fortran
end module fortranmodule
"""

with open("fortranModule.f90", 'w') as fortranfile:
    for line in fsource:
        fortranfile.write(line)

import subprocess
try:
    data = subprocess.check_output(["f2py",
#     data = subprocess.check_output(["/home/pythonstudent/.local/bin/f2py",
                                     "-c", "fortranModule.f90",
                                     "-m", "myfortranmodule",
                                     "--opt='-Ofast -march=native -fopenmp -fvect-cost-m",
                                     #"--opt='-Ofast -xHost -qopenmp '", "--noarch", "-l",
                                     #"--debug-capi", "--debug",
                                     "-DF2PY_REPORT_ON_ARRAY_COPY=1"
                                     ], stderr=subprocess.STDOUT).decode('utf-8')
except subprocess.CalledProcessError as e:
    print(e.output.decode('utf-8'))
else:

```

```

    #print(data)
    print("compilation OK")

```

compilation OK

```

In [41]: import os, sys
        from myfortranmodule import fortranmodule

        _ffunction1 = fortranmodule.function1
        ffunction2 = fortranmodule.function2
        ffunction2.__name__ = "ffunction2"
        convolve_fortran = fortranmodule.convolve_fortran

        def ffunction1(a ,b):
            return _ffunction1(a, b).astype(bool)

        def fortran_convolve(f, g):
            h, err = convolve_fortran(f, g)
            if err:
                print(err)
                raise ValueError("FORTRAN ERROR ! (Probably : Only odd dimensions on filter sup
            return h

        def fortran_setup(f, g):
            # memory ordering for fortran
            ft = np.asfortranarray(f.copy())
            gt = np.asfortranarray(g.copy())
            return ft, gt

```

```

In [42]: ### The produced binary is simple enough to be decompiled
        import subprocess
        data = subprocess.check_output(["gdb", "-batch",
                                         "-ex", "file myfortranmodule.cpython-35m-x86_64-linux-g
                                         "-ex", "disassemble /s __fortranmodule_MOD_function1"]
                                         #"-ex", "disassemble /m fortranmodule_mp_function1_").

        print(data)

```

Dump of assembler code for function __fortranmodule_MOD_function1:
fortranModule.f90:

```

7          subroutine function1(a, b, c, n)
    0x00000000000005cd0 <+0>:      sub    $0x38,%rsp
    0x00000000000005cd4 <+4>:      mov     (%rcx),%rax
    0x00000000000005cd7 <+7>:      xor     %ecx,%ecx

8          implicit none
9          integer(kind=8), intent(in) :: n
10         double precision,intent(in) :: a(n)
11         double precision,intent(in) :: b(n)

```



```

12
13         logical,intent(out)           :: c(n)
14
15         !$OMP PARALLEL WORKSHARE NUM_THREADS(8)
16         0x0000000000005cd9 <+9>:      mov     %rdi, (%rsp)
17         0x0000000000005cdd <+13>:     lea     -0x504(%rip),%rdi      # 0x57e0 <__fortranmodule_MO
18         0x0000000000005ce4 <+20>:     mov     %rsi, 0x8(%rsp)
19         0x0000000000005ce9 <+25>:     mov     %rdx, 0x10(%rsp)
20         0x0000000000005cee <+30>:     mov     %rsp,%rsi
21         0x0000000000005cf1 <+33>:     mov     $0x8,%edx
22         0x0000000000005cf6 <+38>:     mov     %rax, 0x28(%rsp)
23         0x0000000000005cfb <+43>:     mov     %rax, 0x20(%rsp)
24         0x0000000000005d00 <+48>:     mov     %rax, 0x18(%rsp)
25         0x0000000000005d05 <+53>:     callq   0x1f50 <GOMP_parallel@plt>
26
27         c = a * b - 4.1 * a > 2.5 * b
28         !$OMP END PARALLEL WORKSHARE
29
30         end subroutine function1
31         0x0000000000005d0a <+58>:     add     $0x38,%rsp
32         0x0000000000005d0e <+62>:     retq
33
End of assembler dump.

```

```

In [43]: _ = check(ffunction1, (a, b), ref=ref1)
         _ = check(ffunction2, (a, b), ref=ref2)
         print()
         _ = check(fortran_convolve, (small_f, g), ref=ref_small_convolve, setup=fortran_setup)
         _ = check(fortran_convolve, (large_f, g), ref=ref_large_convolve, setup=fortran_setup)

```

```

Testing ffunction1
1.23 ms ± 12 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)
Testing ffunction2
11.5 ms ± 115 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

```

```

Testing fortran_convolve
566 µs ± 4.39 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)
Testing fortran_convolve
59.7 ms ± 1.15 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)

```

- parallelism
- cuda

```

In [44]: """
         CUDA DOESN'T WORK ON THE VIRTUAL MACHINE
         YOU ARE WELCOME TO TRY THIS ON YOU OWN COMPUTER
         """

```

```

from string import Template

cuda_src_template = Template("""
// Cuda splitting
#define MTB ${max_threads_per_block}
#define MBP ${max_blocks_per_grid}

// Array size
#define fx ${fx}
#define fy ${fy}
#define gx ${gx}
#define gy ${gy}

// Macro for converting subscripts to linear index:
#define f_INDEX(i, j) (i) * (fy) + (j)

// Macro for converting subscripts to linear index:
#define g_INDEX(i, j) (i) * (gy) + (j)

__global__ void convolve_cuda(long *f, long *g, long *h) {

    unsigned int idx = blockIdx.y * MTB * MBP + blockIdx.x * MTB + threadIdx.x;

    // Convert the linear index to subscripts:
    unsigned int i = idx / fy;
    unsigned int j = idx % fy;

    long smax = gx;
    long tmax = gy;

    long smid = smax / 2;
    long tmid = tmax / 2;

    if (smid <= i && i < fx - smid) {
        if (tmid <= j && j < fy - tmid) {

            h[f_INDEX(i, j)] = 0.;

            for (long s = 0; s < smax; s++)
                for (long t = 0; t < tmax; t++)
                    h[f_INDEX(i, j)] += g[g_INDEX(s, t)] * f[f_INDEX(i + s - smid, j + t -
        }
    }
}
""")

```

In [45]: """

```

CUDA DOESN'T WORK ON THE VIRTUAL MACHINE
YOU ARE WELCOME TO TRY THIS ON YOU OWN COMPUTER
"""

```

```

import skcuda.misc as misc
import pycuda.autoinit
device = pycuda.autoinit.device
max_threads_per_block, _, max_grid_dim = misc.get_dev_attrs(device)
max_blocks_per_grid = max(max_grid_dim)

```

Error Traceback (most recent call last)

```

/Data/WORK/Formations/Python/formation_python/04.advanced.ipynb in <module>()
4 """
5 import skcuda.misc as misc
----> 6 import pycuda.autoinit
7 device = pycuda.autoinit.device
8 max_threads_per_block, _, max_grid_dim = misc.get_dev_attrs(device)

```

```

/Data/WORK/Formations/Python/formation_python/venv/lib/python3.5/site-packages/pycuda/au
3
4 # Initialize CUDA
----> 5 cuda.init()
6
7 from pycuda.tools import make_default_context

```

Error: cuInit failed: unknown error

In [46]: """

```

CUDA DOESN'T WORK ON THE VIRTUAL MACHINE
YOU ARE WELCOME TO TRY THIS ON YOU OWN COMPUTER
"""

```

```

from functools import partial
from pycuda.compiler import SourceModule

```

```

cuda_src = cuda_src_template.substitute(max_threads_per_block=max_threads_per_block,
                                         max_blocks_per_grid=max_blocks_per_grid,
                                         fx=large_f.shape[0], fy=large_f.shape[1],
                                         gx=g.shape[0], gy=g.shape[1]
                                         )

```

```

cuda_module = SourceModule(cuda_src, options= ["-O3", "-use_fast_math", "-default-stream"]
print("Compilation OK")

```

```

__convolve_cuda = cuda_module.get_function('convolve_cuda')

block_dim, grid_dim = misc.select_block_grid_sizes(device, f.shape)
_convolve_cuda = partial(__convolve_cuda,
                          block=block_dim,
                          grid=grid_dim)

```

NameError Traceback (most recent call last)

```

/Data/WORK/Formations/Python/formation_python/04.advanced.ipynb in <module>()
      6 from pycuda.compiler import SourceModule
      7
----> 8 cuda_src = cuda_src_template.substitute(max_threads_per_block=max_threads_per_block,
      9                                         max_blocks_per_grid=max_blocks_per_grid,
     10                                         fx=large_f.shape[0], fy=large_f.shape[1],

```

NameError: name 'max_threads_per_block' is not defined

```

In [47]: """
        CUDA DOESN'T WORK ON THE VIRTUAL MACHINE
        YOU ARE WELCOME TO TRY THIS ON YOU OWN COMPUTER
        """

import pycuda.gpuarray as gpuarray

def convolve_cuda(f, g):
    f_gpu, g_gpu = cuda_setup(f, g)
    h_gpu = convolve_cuda2(f_gpu, g_gpu)
    return cuda_finish(h_gpu)

def convolve_cuda2(f_gpu, g_gpu):
    h_gpu = gpuarray.zeros_like(f_gpu)
    _convolve_cuda(f_gpu, g_gpu, h_gpu)
    return h_gpu

def cuda_setup(f, g):
    f_gpu = gpuarray.to_gpu(f)
    g_gpu = gpuarray.to_gpu(g)
    return f_gpu, g_gpu

def cuda_finish(h_gpu):
    return h_gpu.get()

```

```

In [48]: """
        CUDA DOESN'T WORK ON THE VIRTUAL MACHINE

```

```
YOU ARE WELCOME TO TRY THIS ON YOU OWN COMPUTER
"""
```

```
print("Cuda")
check(convolve_cuda, (f, g), ref=ref)
print("Cuda without comm")
check(convolve_cuda2, (f, g), ref=ref, setup=cuda_setup, finish=cuda_finish)
print("Finished")
```

Cuda

Testing convolve_cuda

```
-----
LogicError                                Traceback (most recent call last)
```

```
/Data/WORK/Formations/Python/formation_python/04.advanced.ipynb in <module>()
4 """
5 print("Cuda")
----> 6 check(convolve_cuda, (f, g), ref=ref)
7 print("Cuda without comm")
8 check(convolve_cuda2, (f, g), ref=ref, setup=cuda_setup, finish=cuda_finish)

/Data/WORK/Formations/Python/formation_python/04.advanced.ipynb in check(fun, data, ref,
12
13     if ref is None:
---> 14         res = finish(fun(f, g))
15     else:
16         res = finish(fun(f, g))

/Data/WORK/Formations/Python/formation_python/04.advanced.ipynb in convolve_cuda(f, g)
6
7 def convolve_cuda(f, g):
----> 8     f_gpu, g_gpu = cuda_setup(f, g)
9     h_gpu = convolve_cuda2(f_gpu, g_gpu)
10     return cuda_finish(h_gpu)

/Data/WORK/Formations/Python/formation_python/04.advanced.ipynb in cuda_setup(f, g)
16
17 def cuda_setup(f, g):
---> 18     f_gpu = gpuarray.to_gpu(f)
19     g_gpu = gpuarray.to_gpu(g)
20     return f_gpu, g_gpu
```

```

/Data/WORK/Formations/Python/formation_python/venv/lib/python3.5/site-packages/pycuda/gp
990 def to_gpu(ary, allocator=drv.mem_alloc):
991     """converts a numpy array to a GPUArray"""
--> 992     result = GPUArray(ary.shape, ary.dtype, allocator, strides=_compact_strides(ary))
993     result.set(ary)
994     return result

```

```

/Data/WORK/Formations/Python/formation_python/venv/lib/python3.5/site-packages/pycuda/gp
208         if gpudata is None:
209             if self.size:
--> 210                 self.gpudata = self.allocator(self.size * self.dtype.itemsize)
211             else:
212                 self.gpudata = None

```

LogicError: cuMemAlloc failed: initialization error

- asyncio
 - not a real parallelism
 - effective for io-bound tasks (web)
 - not very interesting here
- multithreading
 - more parallelism (GIL)
 - shared memory
 - two main implementation
 - * threading (stdlib) which is flexible
 - * joblib which is relatively easy to use
- multiprocessing
 - real parallelism
 - limited to one computer
 - two main implementation
 - * multiprocessing (stdlib) which is flexible
 - * joblib which is relatively easy to use
- mpi (mpi4py)
 - real parallelism
 - unlimited
 - relatively complex to use (same as in C, fortran, ...)

```

In [49]: import time
import numpy as np
def heavy_fonction(i):

```

```

t = np.random.rand() / 10
time.sleep(t)
return i, t

```

In [50]: `from joblib import Parallel, delayed`

```

if __name__ == "__main__":

    tic = time.time()
    res = Parallel(n_jobs=-1, backend='threading')(delayed(heavy_fonction)(i) \
                                                    for i in range(2000))

    tac = time.time()
    index, times = np.asarray(res).T
    print(tac - tic)
    print(times.sum())

```

```

12.737246036529541
100.36369606025077

```

In [51]: `from threading import Thread, RLock`

```

N = 2000
N_t = 10
current = 0
nprocs = 8
output_list = np.empty(N)

lock = RLock()

class ThreadJob(Thread):
    def run(self):
        """This code will be executed by each thread"""
        global current

        while current < N:

            with lock:
                position = current
                current += N_t

            fin = min(position + N_t + 1, N)

            for i in range(position, fin):
                j, t = heavy_fonction(i)
                output_list[j] = t

```

```

if __name__ == "__main__":

    # Threads creation
    threads = [ThreadJob() for i in range(nprocs)]

    tic = time.time()
    # Threads starts
    for thread in threads:
        thread.start()

    # Waiting that all thread have finish
    for thread in threads:
        thread.join()
    tac = time.time()

    print(tac - tic)
    print(output_list.sum())

```

```

14.059715747833252
99.97539302441548

```

```

In [52]: import multiprocessing as mp
         from queue import Empty

         def process_job(q,r):
             """This code will be executed by each process"""
             while True:
                 try:
                     i = q.get(block=False)
                     r.put(heavy_fonction(i))
                 except Empty:
                     if q.empty():
                         if q.qsize() == 0:
                             break

         if __name__ == "__main__":

             # Define an output queue
             r = mp.Queue()

             # Define an input queue
             q = mp.Queue()

             for i in range(2000):
                 q.put(i)

```



```

nprocs = 8
# Setup a list of processes that we want to run
processes = [mp.Process(target=process_job, args=(q, r)) for i in range(nprocs)]

tic = time.time()

# Run processes
for p in processes:
    p.start()

# Get process results from the output queue
results = np.empty(2000)
for i in range(2000):
    j, t = r.get()
    results[j] = t

tac = time.time()

# Exit the completed processes
for p in processes:
    p.join()

print(tac - tic)
print(results.sum())

```

12.583194017410278

99.33507414916964

2.6.1 Exercise

The following code read an image in pgm format (ascii) and store it in a 2D list.

For each pixel of the image a kernel get all neighbors (9 counting the pixel itself) and apply a computation.

Analyse the performance of the code, identify bottleneck and try to optimize it.

In [53]: %matplotlib notebook

```

import math
import matplotlib.pyplot as plt

def get_description(filename):
    """
    Read the header part of the file
    """
    f = open(filename, 'r')
    nline = 0
    description = {}

```

```

while nline < 3:
    line = f.readline()
    if line[0] == '#':
        continue
    nline += 1
    if nline == 1:
        description['format'] = line.strip()
    elif nline == 2:
        description['dimension'] = int(line.split()[1]), int(line.split()[0])
    elif nline == 3:
        description['deep'] = int(line.strip())
f.close()
return description

def get_value(filename, coord):
    """
    Get value at coord in an image in the PGM format

    The main problem here is that the file have a limited width, and the values are wrapped.
    Thus, the value at coord 12,32 might be in the 24,6 in the file
    """
    description = get_description(filename)
    xdim, ydim = description['dimension']
    i = coord[0]
    j = coord[1]
    f = open(filename, 'r', encoding='utf-8')
    nline = 0
    while nline < 3:
        line = f.readline()
        if line[0] == '#':
            continue
        nline += 1
    #here we are at coordinate (0,0)
    icur, jcur = 0,0
    test = True
    while(test):
        values = f.readline().split()
        nvalues = len(values)
        if (icur == i):
            if (jcur + nvalues > j):
                jvalues = j - jcur
                value = values[jvalues]
                test=False
            else:
                jcur += nvalues
        else:
            jcur += nvalues
        if (jcur >= ydim):

```

```

        icur += jcur // ydim
        jcur = jcur % ydim
    f.close()
    return int(value)

def read_file(filename):
    """
    Read an image in the PGM format
    """
    description=get_description(filename)
    values = []
    for i in range(description['dimension'][0]):
        values.append([])
        for j in range(description['dimension'][1]):
            values[i].append(get_value(filename, (i, j)))

    return values

def get_neighbors(tab, i, j):
    neigh = []
    for jrange in [-1, 0, 1]:
        for irange in [-1, 0, 1]:
            neigh.append(tab[i + irange][j + jrange])
    return neigh

def compute_wtf(neigh):
    value = 1.
    for i in range(len(neigh)):
        value *= math.exp(neigh[i]) ** (1 / len(neigh))
    value = math.log(value)

    return float(value)

def kernel(tab):
    """
    Apply compute_wtf on each pixel except boundary
    """
    xdim = len(tab)
    ydim = len(tab[0])
    result = []
    #1st line
    result.append([0])
    for jrange in range(1, ydim):
        result[0].append(0)
    for irange in range(1, xdim - 1):
        #1st column
        result.append([0])

```

```

        for jrange in range(1, ydim - 1):
            neigh = get_neighbors(tab, irange, jrange)
            result[irange].append(compute_wtf(neigh))
        #last column
        result[irange].append(0)
    #last line
    result.append([])
    for jrange in range(ydim):
        result[xdim - 1].append(0)
    return result

def job(data):
    """
    Apply kernel of each image
    """
    results = []
    for image in data:
        results.append(kernel(image))
    return results

def init(files):
    """
    Read all files
    """
    data = []
    for file in files:
        data.append(read_file(file))

    return data

def plot(data):
    nimages = len(data)

    if nimages > 1:
        fig, axes = plt.subplots(nimages, 1)
        for image, ax in zip(data, axes):
            ax.imshow(image)
    else:
        plt.figure()
        plt.imshow(data[0])

    plt.show()

files=["data/test.pgm"]

if __name__ == "__main__":

```

```

print("Reading Files")
data = init(files)

print("Computing")
result = job(data)

plot(data)
plot(result)

```

Reading Files
Computing

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

Try to optimize this code.

You can apply your code on the following images : - data/test.pgm - data/test32.pgm - data/brain_604.ascii.pgm - data/apollonian_gasket.ascii.pgm - data/dla.ascii.pgm

For reference, the timing on my computer are :

For data/test.pgm

On my computer :

Reading Files

6.67 ms ± 262 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

Computing

503 µs ± 5.41 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)

My solution

Reading Files

83.9 µs ± 3.9 µs per loop (mean ± std. dev. of 7 runs, 10000 loops each)

Computing

255 µs ± 19.4 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)

And, the bigger the image, the bigger the gain !

2.6.2 Hints

Part 1 - Open input file only once - Avoid appending data - Use numpy array for data storage - What is really doing the compute_wtf function ?

Part 2 - compile the compute part

Part 3 - parallelize the work on each image

2.6.3 Solution part 1

```
import numpy as np
from copy import deepcopy

def get_description(file):
    """
    Read the header part of the file

    And leave the file at the end of header (start of values)
    """
    # return to beginning
    file.seek(0)
    nline = 0
    description = {}
    while nline < 3:
        line = file.readline()
        if line[0] == '#':
            continue
        nline += 1
        if nline == 1:
            description['format'] = line.strip()
        elif nline == 2:
            description['dimension'] = int(line.split()[1]), int(line.split()[0])
        elif nline == 3:
            description['deep'] = int(line.strip())
    return description

def read_values(file, description):
    """
    Read all the values directly
    """
    # pre-allocate the array
    nx, ny = description['dimension']
    values = np.empty((nx * ny))
    i = 0
    for line in file:
        if line[0] == '#':
            continue
        vals = line.split()
        nvals = len(vals)
        values[i:i + nvals] = [int(v) for v in vals]
        i += nvals
    return values.reshape((nx, ny))

def read_file(filename):
    """
    Read an image in the PGM format
    """
```

```

    """
    # open the file once
    with open(filename, 'r', encoding="utf-8") as file:

        # read the header part
        description = get_description(file)

        # read the values
        values = read_values(file, description)
    return values

def init(files):
    """
    Read all files
    """
    data = []
    for file in files:
        data.append(read_file(file))

    return data

#prepare result array
result = deepcopy(data)

```

2.6.4 Solution part 2

In the python part :

```

def kernel(i):
    """
    Apply compute_wtf on each pixel except boundary
    """
    global data, result, t_omp
    result[i] = ckernel(data[i], nt_omp)

```

In a new cell : %load_ext Cython

In a new cell :

```

%%cython --compile-args=-fopenmp --link-args=-fopenmp
## --compile-args=-DCYTHON_TRACE_NOGIL=1 --compile-args=-DCYTHON_TRACE=1
# cython: language_level=3
# cython: boundscheck=False
# cython: wraparound=False
# cython: initializedcheck=False
# cython: binding=True
import numpy as np

```

```

cimport numpy as np
cimport cython
from cython.parallel cimport parallel, prange

def ckernel(double[:,::1] data, long nt):
    cdef long n = data.shape[0]
    cdef long m = data.shape[1]

    cdef double[:,::1] res = np.zeros([n, m], dtype=np.double)
    cdef double value

    cdef long i, j, s, t
    with nogil, parallel(num_threads=nt):
        for i in prange(1, n - 1):
            for j in range(1, m - 1):
                value = 0
                for s in range(-1, 2):
                    for t in range(-1, 2):
                        value += data[i + s, j + t]
                res[i, j] += value / 9
    return res

```

2.6.5 Solution part 3

```

from threading import Thread, RLock
import os

nprocs = os.cpu_count()
nt_omp = nprocs // 2
nt_job = nprocs - nt_omp

result = []
data = []

current = 0

verrou = RLock()

class ThreadJob(Thread):
    def run(self):
        global current, verrou
        """Code à exécuter pendant l'exécution du thread."""
        while current < len(data):

            with verrou:
                position = current
                current += 1

```



```

        kernel(position)

def job(data):
    """
    Apply kernel on each image
    """
    global current

    current = 0
    # Création des threads
    threads = [ThreadJob() for i in range(nt_job)]

    # Lancement des threads
    for thread in threads:
        thread.start()

    # Attend que les threads se terminent
    for thread in threads:
        thread.join()

#sort data bigger first for better equilibrium
data = sorted(data, key=np.size, reverse=True)

```

2.6.6 Full Solution

Cell 1: %load_ext Cython

Cell 2:

```

%%cython --compile-args=-fopenmp --link-args=-fopenmp
## --compile-args=-DCYTHON_TRACE_NOGIL=1 --compile-args=-DCYTHON_TRACE=1
# cython: language_level=3
# cython: boundscheck=False
# cython: wraparound=False
# cython: initializedcheck=False
# cython: binding=True
import numpy as np
cimport numpy as np
cimport cython
from cython.parallel cimport parallel, prange

def ckernel(double[:,::1] data, long nt):
    cdef long n = data.shape[0]
    cdef long m = data.shape[1]

    cdef double[:,::1] res = np.zeros([n, m], dtype=np.double)
    cdef double value

```

```

cdef long i, j, s, t
with nogil, parallel(num_threads=nt):
    for i in prange(1, n - 1):
        for j in range(1, m - 1):
            value = 0
            for s in range(-1, 2):
                for t in range(-1, 2):
                    value += data[i + s, j + t]
            res[i, j] += value / 9
return res

```

Cell 3 :

```

%matplotlib notebook

import numpy as np
import matplotlib.pyplot as plt
from threading import Thread, RLock
from copy import deepcopy
import os

nprocs = os.cpu_count()
nt_omp = nprocs // 2
nt_job = nprocs - nt_omp

result = []
data = []

current = 0

verrou = RLock()

class ThreadJob(Thread):
    def run(self):
        global current, verrou
        """Code à exécuter pendant l'exécution du thread."""
        while current < len(data):

            with verrou:
                position = current
                current += 1

            kernel(position)

def get_description(file):
    """
    Read the header part of the file

```

```

And leave the file at the end of header (start of values)
"""

# return to beginning
file.seek(0)
nline = 0
description = {}
while nline < 3:
    line = file.readline()
    if line[0] == '#':
        continue
    nline += 1
    if nline == 1:
        description['format']=line.strip()
    elif nline == 2:
        description['dimension']=int(line.split()[1]), int(line.split()[0])
    elif nline == 3:
        description['deep']=int(line.strip())
return description

def read_values(file, description):
    """
    Read all the values directly
    """
    # pre-allocate the array
    nx, ny = description['dimension']
    values = np.empty((nx * ny))
    i = 0
    for line in file:
        if line[0] == '#':
            continue
        vals = line.split()
        nvals = len(vals)
        values[i:i + nvals] = [int(v) for v in vals]
        i += nvals
    return values.reshape((nx, ny))

def read_file(filename):
    """
    Read an image in the PGM format
    """
    # open the file once
    with open(filename, 'r', encoding="utf-8") as file:

        # read the header part
        description = get_description(file)

        # read the values

```

```

        values = read_values(file, description)
    return values

def kernel(i):
    """
    Apply compute_wtf on each pixel except boundary
    """
    global data, result, t_omp
    result[i] = ckernel(data[i], nt_omp)

def job(data):
    """
    Apply kernel of each image
    """
    global current

    current = 0
    # Création des threads
    threads = [ThreadJob() for i in range(nt_job)]

    # Lancement des threads
    for thread in threads:
        thread.start()

    # Attend que les threads se terminent
    for thread in threads:
        thread.join()

def init(files):
    """
    Read all files
    """
    data = []
    for file in files:
        data.append(read_file(file))

    return data

def plot(data):
    nimages = len(data)
    if nimages > 1:
        fig, axes = plt.subplots(nimages, 1)
        for image, ax in zip(data, axes):
            ax.imshow(image)
    else:
        plt.figure()
        plt.imshow(data[0])

```

```

plt.show()

files = ["data/test.pgm",
         "data/test32.pgm",
         "data/brain_604.ascii.pgm",
         "data/apolonian_gasket.ascii.pgm",
         "data/dla.ascii.pgm",
        ]

if __name__ == "__main__":
    data = init(files)

    #sort data bigger first for better equilibrium
    data = sorted(data, key=np.size, reverse=True)

    #prepare result array
    result = deepcopy(data)

    #plot(data)
    plot(result)

```